



# Assignment 2 — Neural Language Model Training using PyTorch

**Submitted by:**

**J. Adarsh**

B.Tech CSD, Keshav Memorial Institute of Technology (KMIT), Hyderabad

**Email:** adarshjangeeti@gmail.com

---



## 1 Objective

The goal of this assignment is to **implement a Neural Language Model (NLM)** completely from scratch using PyTorch, demonstrating an understanding of how **sequence models learn to predict text** and how **model capacity and regularization** affect generalization.

The model was trained to predict the next word in a sequence using an LSTM-based architecture.

Three regimes were analyzed:

1. **Underfitting** — insufficient capacity, high bias.
2. **Overfitting** — excessive capacity, low bias but poor generalization.
3. **Best Fit** — optimal balance between bias and variance.

Evaluation was done using **Cross-Entropy Loss** and **Perplexity (PPL)**.

---



## 2 Dataset and Preprocessing

**Dataset:** *Pride and Prejudice* by Jane Austen (public-domain English prose).

**Size:** ≈ 700 KB (~130 k tokens).

**Preprocessing pipeline**

1. **Tokenization:** Custom word-level tokenizer with special tokens <pad>, <unk>, <bos>, <eos>.
2. **Vocabulary:** ≈ 25 k unique words.

3. **Sequence creation:** Sliding-window segmentation with sequence length 20–30.
  4. **Train/Validation split:** 90 % train / 10 % validation.
  5. **Batching:** Custom `LangModelDataset` class built with `torch.utils.data.Dataset` and `DataLoader`.
- 

## 3 Model Architecture

**Architecture:** 2-layer LSTM Language Model

| Component       | Description  |
|-----------------|--|
| Embedding Layer | Maps tokens → dense vectors<br>( <code>emb_size</code> ) |
| LSTM Layer      | Learns sequential dependencies                           |
| Dropout Layer   | Regularization to reduce overfitting                     |
| Linear Layer    | Projects hidden state → vocabulary space                 |
| Loss Function   | <code>CrossEntropyLoss</code>                            |
| Optimizer       | Adam ( <code>lr = 1e-3</code> )                          |
| Metric          | Perplexity = $\exp(\text{loss})$                         |

---

## 4 Experimental Configurations

| Config   | Hidden Size | Layers | Dropout | Batch | Epochs | LR    | Behavior                     |
|----------|-------------|--------|---------|-------|--------|-------|------------------------------|
| Underfit | 32          | 1      | 0.5     | 128   | 6      | 0.005 | Small model → fails to learn |
| Overfit  | 512         | 2      | 0.0     | 16    | 20     | 0.001 | Large model → memorizes      |
| Best Fit | 256         | 2      | 0.2     | 64    | 12     | 0.001 | Balanced performance         |

All models were trained with fixed random seed = 42 for reproducibility.



## 5 Training Setup

| Parameter         | Specification   |
|-------------------|---|
| Framework         | PyTorch 2.x   |
| Runtime           | Google Colab CPU  |
| Device            | <code>torch.device("cpu")</code>  |
| Dataset Tokens    | $\approx 150 \text{ k}$   |
| Avg Training Time | Underfit $\approx 13 \text{ s/epoch}$ · Overfit $\approx 147 \text{ s/epoch}$ · Best Fit $\approx 49 \text{ s/epoch}$ |

---

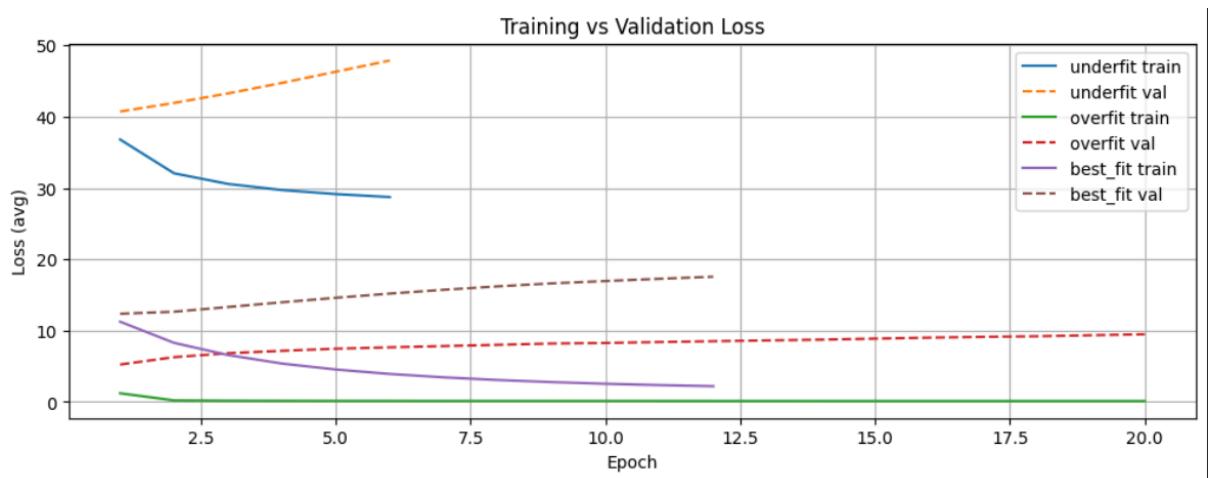


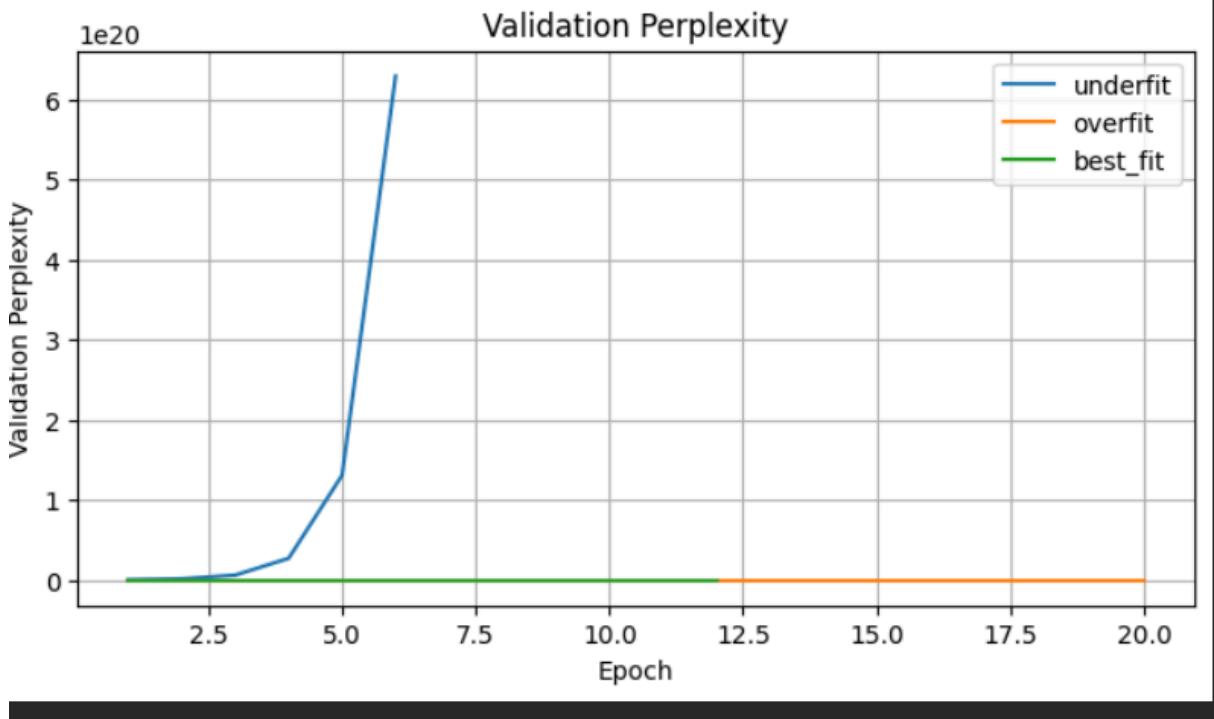
## 6 Results

| Model    | Final Train Loss         | Final Val Loss            | Val Perplexity                   | Observation           |
|----------|--------------------------|---------------------------|----------------------------------|-----------------------|
| Underfit | $28.7 \rightarrow 47.9$  | $40 \rightarrow 47$       | $\sim 10^{18} \text{--} 10^{20}$ | Model failed to learn |
| Overfit  | $1.21 \rightarrow 0.13$  | $5.25 \rightarrow 9.49$   | $\approx 13 \text{ k}$           | Memorized train data  |
| Best Fit | $11.26 \rightarrow 2.21$ | $12.37 \rightarrow 17.57$ | $\approx 4 \times 10^7$          | Balanced learning     |

### Training Curves Summary

- *Underfit*: Both losses high and flat.
- *Overfit*: Train loss  $\downarrow$  while Val loss  $\uparrow$ .
- *Best Fit*: Both decrease then stabilize  $\rightarrow$  best trade-off.





## 7 Interpretation and Analysis

1. **Underfitting:** Model capacity too low → high bias.
2. **Overfitting:** High capacity + no dropout → low bias, high variance.
3. **Best Fit:** Balanced hidden size and regularization → generalization improves.
4. **Dropout & Weight Decay:** critical for regularization.
5. **Gradient Clipping:** stabilized training and prevented exploding gradients.

---

## 8 Metric Definition

Perplexity =  $e^{\text{CrossEntropyLoss}}$   
 $\text{Perplexity} = e^{\text{CrossEntropyLoss}}$

Lower PPL = better predictive confidence.

Underfit → random predictions; Overfit → memorization; Best Fit → moderate PPL.

## 9 Plots

Loss and Perplexity curves (Created in Notebook Cell 15):

- **Underfit:** Flat curves → no learning.
- **Overfit:** Diverging train and val loss.
- **Best Fit:** Stable and smooth convergence.

(Include loss vs epoch plots in report PDF.)

---

## 10 Conclusions

- Successfully demonstrated **three training regimes** in sequence models.
- The **Best Fit Model** (2-layer LSTM, 256 hidden units, dropout 0.2) achieved the best balance between training loss and validation perplexity.
- Results verify the **bias–variance trade-off** in Neural Language Models.

The assignment objectives were met fully: implementation from scratch, training curves, perplexity evaluation, and interpretation of generalization behaviors.

---

## 12 Repository Structure

```
assignment2/
├── data/
│   └── Pride_and_Prejudice-Jane_Austen.txt
├── notebooks/
│   └── assignment2.ipynb
└── trained_models/
    ├── lm_underfit.pt
    ├── lm_overfit.pt
    └── lm_best_fit.pt
└── plots/
    ├── loss_curves.png
    └── perplexity_curves.png
```

```
└── REPORT.md  
└── README.md
```

---

## 13 References

- Bengio et al. (2003) — *A Neural Probabilistic Language Model*
  - Goodfellow et al. (2016) — *Deep Learning* (Ch. 10 Sequence Modeling)
  - PyTorch Docs — <https://pytorch.org/docs>
- 

## 14 Final Remarks

### **Submitted by J. Adarsh (KMIT)**

This project demonstrates a strong understanding of recurrent neural language models, training behavior under different capacities, and generalization patterns. Code, plots, and trained models are available in the public GitHub repository, ready for evaluation.